

Virtual time & Global states

Thoai Nam


High Performance Computing Lab (HPC Lab)
Faculty of Computer Science and Engineering
HCMC University of Technology

Contents

- Time ordering and clock synchronization
- Virtual time (logical clock)
- Distributed snapshot (global state)
- Consistent/Inconsistent global state
- Rollback Recovery
- Debugging & Race messages

Clock synchronization

- Time is unambiguous in centralized systems
 - System clock keeps time, all entities use this for time
- Distributed systems: each node has own system clock
 - Crystal-based clocks are less accurate (1 part in million)
 - *Problem:* An event that occurred after another may be assigned an earlier time




Physical clocks: a primer

- Accurate clocks are atomic oscillators
 - 1s ~ 9,192,631,770 transitions of the cesium 133 atom
- Most clocks are less accurate (e.g., mechanical watches)
 - Computers use crystal-based blocks (one part in million)
 - Results in *clock drift*
- How do you tell time?
 - Use astronomical metrics (solar day)
- Universal coordinated time (*UTC*) – international standard based on atomic time
 - Add leap seconds to be consistent with astronomical time
 - UTC broadcast on radio (satellite and earth)
 - Receivers accurate to 0.1 – 10 ms
- Need to synchronize machines with a master or with one another


Clock synchronization

- Each clock has a maximum drift rate r
 - $1-r \leq dC/dt \leq 1+r$
 - Two clocks may drift by $2r \cdot D_t$ in time D_t
 - To limit drift to $d \rightarrow$ resynchronize every $d/2r$ seconds



Cristian's algorithm


- Synchronize machines to a *time server* with a UTC receiver
- Machine P requests time from server every $d/2r$ seconds
 - Receives time t from server, P sets clock to $t + t_{reply}$ where t_{reply} is the time to send reply to P
 - Use $(t_{req} + t_{reply})/2$ as an estimate of t_{reply}
 - Improve accuracy by making a series of measurements



Berkeley algorithm

- Used in systems without UTC receiver
 - Keep clocks synchronized with one another
 - One computer is *master*, other are *slaves*
 - Master periodically polls slaves for their times
 - ✧ Average times and return differences to slaves
 - ✧ Communication delays compensated as in Cristian's algorithm
 - Failure of master => election of a new master

Berkeley algorithm



- a) The time daemon asks all the other machines for their clock values
- b) The machines answer
- c) The time daemon tells everyone how to adjust their clock

Distributed approaches

- Both approaches studied thus far are centralized
- Decentralized algorithms: use resynchronization intervals
 - Broadcast time at the start of the interval
 - **Collect** all other broadcast that arrive in a period S
 - Use average value of all reported times
 - Can throw away few highest and lowest values
- Approaches in use today
 - *rdate*: synchronizes a machine with a specified machine
 - Network Time Protocol (NTP)
 - ✧ Uses advanced techniques for accuracies of 1-50 ms

Logical clocks

- For many problems, internal consistency of clocks is important
 - Absolute time is less important
 - Use *logical* clocks
- Key idea:
 - Clock synchronization need not be absolute
 - If two machines do not interact, no need to synchronize them
 - More importantly, processes need to agree on the *order* in which events occur rather than the *time* at which they occurred.

Event ordering

- *Problem:* define a total ordering of all events that occur in a system
- Events in a single processor machine are totally ordered
- In a distributed system:
 - No global clock, local clocks may be unsynchronized
 - Cannot order events on different machines using local times
- Key idea [Leslie Lamport]
 - Processes exchange messages
 - Message must be sent before received
 - Send/receive used to order events (and synchronize clocks)


Happened-Before relation

- ❑ If A and B are events in the same process and A executed before B , then $A \rightarrow B$
- ❑ If A represents sending of a message and B is the receipt of this message, then $A \rightarrow B$
- ❑ Relation is transitive: $(A \rightarrow B) \cap (B \rightarrow C) \Rightarrow A \rightarrow C$
- ❑ Relation is undefined across processes that do not exchange messages
 - Partial ordering on events


Event ordering using HB

- Goal: define the notion of time of an event such that
 - If $A \rightarrow B$ then $C(A) < C(B)$
 - If A and B are concurrent, then $C(A) <, = \text{ or } > C(B)$
- Solution:
 - Each processor maintains a logical clock LC_i
 - Whenever an event occurs locally at i , $LC_i = LC_i + 1$
 - When i sends message to j , piggyback LC_i
 - When j receives message from i
 - » If $LC_j < LC_i$ then $LC_j = LC_i + 1$ else do nothing
 - Claim: this algorithm meets the above goals

Lamport's logical clocks



(a)



(b)

More Canonical problems



- Causality
 - Vector timestamps
- Global state and termination detection
- Election algorithms

Causality

- Lamport's logical clocks
 - If $A \rightarrow B$ then $C(A) < C(B)$
 - Reverse is not true!!
 - ✧ Nothing can be said about events by comparing time-stamps!
 - ✧ If $C(A) < C(B)$, then ???
- Need to maintain *causality*
 - Causal delivery: If $\text{send}(m) \rightarrow \text{send}(n) \Rightarrow \text{deliver}(m) \rightarrow \text{deliver}(n)$
 - Capture causal relationships between groups of processes
 - Need a time-stamping mechanism such that:
 - ✧ If $T(A) < T(B)$ then A should have causally preceded B

Vector Clocks




- Each process i maintains a vector V_i
 - $V_i[i]$: number of events that have occurred at process i
 - $V_i[j]$: number of events occurred at process j that process i knows
- Update vector clocks as follows
 - Local event: increment $V_i[i]$
 - Send a message: piggyback entire vector V
 - Receipt of a message:
 - ✧ $V_j[j] = V_j[j] + 1$
 - ✧ Receiver is told about how many events the sender knows occurred at another process k
 $V_j[k] = \max(V_j[k], V_i[k])$
- *Homework*: convince yourself that if $V(A) < V(B)$, then A causally precedes B




Vector Clocks: Happened–Before

Event A in process i has the vector V_A and event B in process j has the vector V_B :

- $A \rightarrow B \Leftrightarrow V_B[i] \geq V_A[i]$
- $B \rightarrow A \Leftrightarrow V_A[j] \geq V_B[j]$
- $A \parallel B \Leftrightarrow \neg(A \rightarrow B) \wedge \neg(B \rightarrow A)$
 $\Leftrightarrow (V_B[i] < V_A[i]) \wedge (V_A[j] < V_B[j])$






$A_1 \rightarrow A_2, A_2 \rightarrow A_3, A_3 \rightarrow A_4$
 $B_1 \rightarrow B_2, B_2 \rightarrow B_3$
 $C_1 \rightarrow C_2, C_2 \rightarrow C_3$

$A_2 \rightarrow B_1, B_2 \rightarrow C_2$
 $B_3 \rightarrow A_3, C_3 \rightarrow A_4$

$A_1 \rightarrow B_1, A_2 \rightarrow B_2, B_1 \rightarrow C_2, B_2 \rightarrow A_3$
 $B_2 \rightarrow C_3, C_2 \rightarrow A_4, B_3 \rightarrow A_4$

$A_1 \rightarrow B_2, A_1 \rightarrow C_2, A_1 \rightarrow C_3, A_1 \rightarrow B_3$
 $B_1 \rightarrow C_3, B_1 \rightarrow A_4, B_3 \rightarrow A_4$
 $C_1 \rightarrow A_4, \dots$

$A_1 \parallel C_1, A_2 \parallel C_1, A_3 \parallel C_1$
 $B_1 \parallel C_1, B_2 \parallel C_1, B_3 \parallel C_1$
 $B_3 \parallel C_2, A_3 \parallel C_2, \dots$



Did you see their relations?


$A_1 \rightarrow A_2, A_2 \rightarrow A_3, A_3 \rightarrow A_4$
 $B_1 \rightarrow B_2, B_2 \rightarrow B_3$
 $C_1 \rightarrow C_2, C_2 \rightarrow C_3$

$A_2 \rightarrow B_1, B_2 \rightarrow C_2$
 $B_3 \rightarrow A_3, C_3 \rightarrow A_4$

$A_1 \rightarrow B_1, A_2 \rightarrow B_2, B_1 \rightarrow C_2, B_2 \rightarrow A_3$
 $B_2 \rightarrow C_3, C_2 \rightarrow A_4, B_3 \rightarrow A_4$

$A_1 \rightarrow B_2, A_1 \rightarrow C_2, A_1 \rightarrow C_3, A_1 \rightarrow B_3$
 $B_1 \rightarrow C_3, B_1 \rightarrow A_4, B_3 \rightarrow A_4$
 $C_1 \rightarrow A_4$

$A_1 \parallel C_1, A_2 \parallel C_1, A_3 \parallel C_1$
 $B_1 \parallel C_1, B_2 \parallel C_1, B_3 \parallel C_1$
 $B_3 \parallel C_2, A_3 \parallel C_2, \dots$



Did you see their relations?

$A_1 \rightarrow A_2, A_2 \rightarrow A_3, A_3 \rightarrow A_4$
 $B_1 \rightarrow B_2, B_2 \rightarrow B_3$
 $C_1 \rightarrow C_2, C_2 \rightarrow C_3$

$A_2 \rightarrow B_1, B_2 \rightarrow C_2$
 $B_3 \rightarrow A_3, C_3 \rightarrow A_4$

$A_1 \rightarrow B_1, A_2 \rightarrow B_2, B_1 \rightarrow C_2, B_2 \rightarrow A_3$
 $B_2 \rightarrow C_3, C_2 \rightarrow A_4, B_3 \rightarrow A_4$


$A_1 \rightarrow B_2, A_1 \rightarrow C_2, A_1 \rightarrow C_3, A_1 \rightarrow B_3$
 $B_1 \rightarrow C_3, B_1 \rightarrow A_4, B_3 \rightarrow A_4$
 $C_1 \rightarrow A_4$

$A_1 \parallel C_1, A_2 \parallel C_1, A_3 \parallel C_1$
 $B_1 \parallel C_1, B_2 \parallel C_1, B_3 \parallel C_1$
 $B_3 \parallel C_2, A_3 \parallel C_2, \dots$


Global State

- Global state of a distributed system
 - Local state of each process
 - Messages sent but not received (state of the queues)
- Many applications need to know the state of the system
 - Failure recovery, distributed deadlock detection
- Problem: how can you figure out the state of a distributed system?
 - Each process is independent
 - No global clock or synchronization
- Distributed snapshot: a consistent global state

Consistent/Inconsistent cuts



(a) A consistent cut




(b) An inconsistent cut

Distributed snapshot algorithm


- Assume each process communicates with another process using unidirectional point-to-point channels (e.g, TCP connections)
- Any process can initiate the algorithm
 - Checkpoint local state
 - Send **marker** on every outgoing channel
- On receiving a marker
 - Checkpoint state if first marker and send marker on outgoing channels, save messages on all other channels until:
 - ✧ Subsequent marker on a channel: stop saving state for that channel

Distributed Snapshot

- A process finishes when
 - It receives a marker on each incoming channel and processes them all
 - State: local state plus state of all channels
 - Send state to initiator
- Any process can initiate snapshot
 - Multiple snapshots may be in progress
 - ✧ Each is separate, and each is distinguished by tagging the marker with the initiator ID (and sequence number)

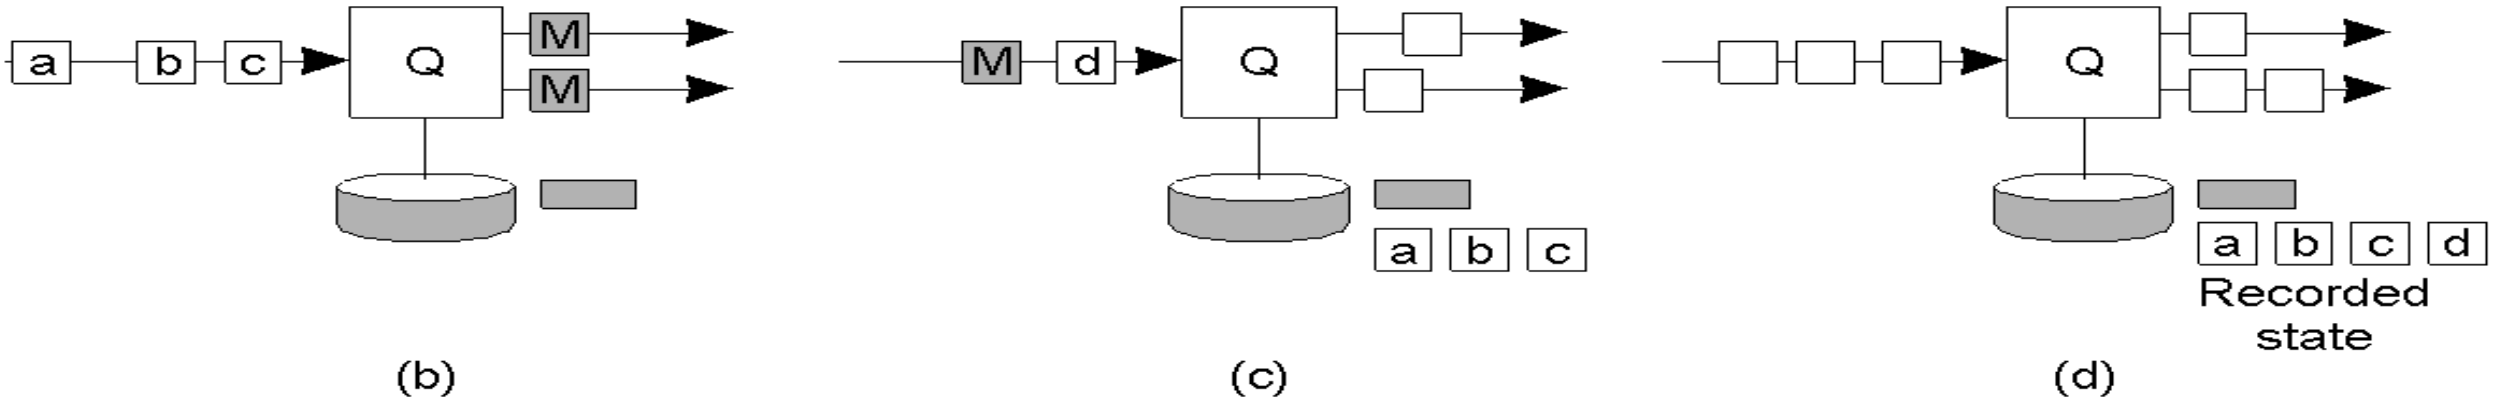


Snapshot algorithm example (I)



(a) Organization of a process and channels for a distributed snapshot

Snapshot algorithm example (2)




- (b) Process **Q** receives a marker for the first time and records its local state
- (c) **Q** records all incoming messages
- (d) **Q** receives a marker from its incoming channel and finishes recording the state of the incoming channel

Recovery

- Techniques thus far allow failure handling
- Recovery: operations that must be performed after a failure to recover to a correct state
- Techniques:
 - Checkpointing:
 - ✧ Periodically checkpoint state
 - ✧ Upon a crash roll back to a previous checkpoint with *a consistent state*

Independent checkpointing



- Each processes periodically checkpoints independently of other processes
- Upon a failure, work backwards to locate a consistent cut
- Problem: if most recent checkpoints form inconsistent cut, will need to keep rolling back until a consistent cut is found
- Cascading rollbacks can lead to a domino effect.

Coordinated checkpointing


- Take a distributed snapshot
- Upon a failure, roll back to the latest snapshot
 - All process restart from the latest snapshot

Message logging

- Checkpointing is expensive
 - All processes restart from previous consistent cut
 - Taking a snapshot is expensive
 - Infrequent snapshots => all computations after previous snapshot will need to be redone [wasteful]
- Combine checkpointing (expensive) with message logging (cheap)
 - Take infrequent checkpoints
 - Log all messages between checkpoints to local stable storage
 - To recover: simply replay messages from previous checkpoint
 - ✧ Avoids recomputations from previous checkpoint


Debugging

Debugging



- Error is detected at B_4
- Where is the root of error?
 - P_1 : Somewhere before B_4
 - P_2, P_3 ?

Debugging





- Error is detected at B_4
- Where is the root of error?
 - P_1 : Somewhere before B_4
 - P_0, P_2 ?

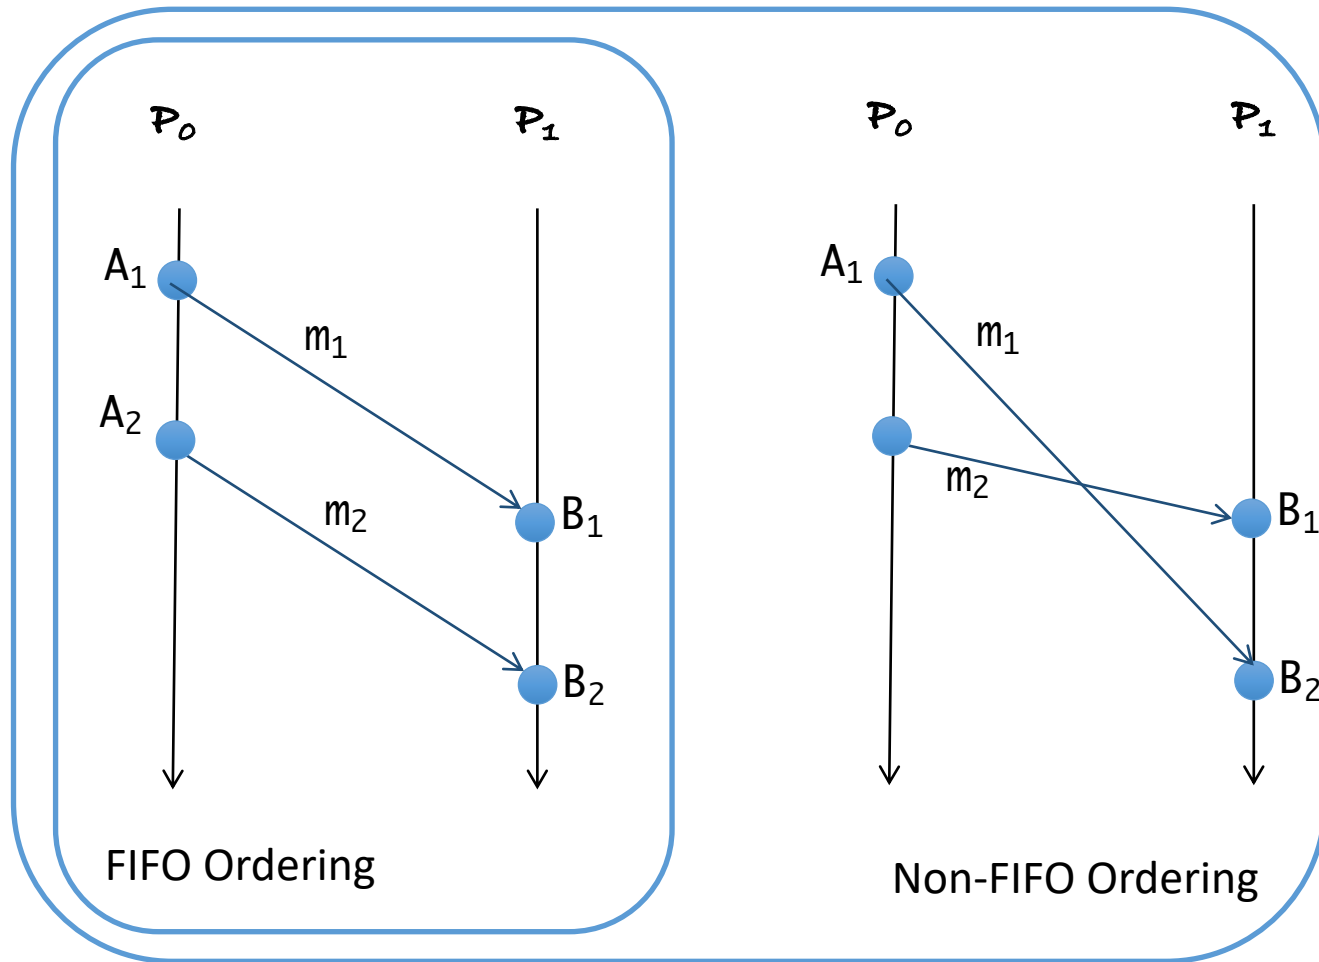
- Event $E \rightarrow B_4$, then actions at E may lead to the error at B_4
- Debugging area = $Past(B_4)$

Race messages

Race messages




Non-determinism of parallel/distributed program



- **FIFO ordering per process**

$(\text{Send}(m_1) \rightarrow \text{Send}(m_2)) \Rightarrow (\text{Receive}(m_1) \rightarrow \text{Receive}(m_2))$

Race messages



- Receive from any source

FIFO ordering per process

- $(B_2 \parallel C_2) \Rightarrow m_4$ may be received at B_2
- The opposite is wrong?
 - $(A_1 \rightarrow B_3)$ but m_1 may be received at B_3
- Finding all race messages at a receive event is a problem