

Distributed Systems

# Distributed File Systems

Thoai Nam

High Performance Computing Lab (HPC Lab)  
Faculty of Computer Science and Engineering  
HCMC University of Technology

# Contents

- Distributed File System architecture
- NFS, HDFS
- ...

# What is a file system?

- Persistent stored data sets
- Hierarchical name space visible to all processes
- API with the following characteristics:
  - access and update operations on persistently stored data sets
  - Sequential access model (with additional random facilities)
- Sharing of data between users, with access control
- Concurrent access:
  - certainly for read-only access
  - what about updates?
- Other features:
  - mountable file stores
  - more? ...

# What is a file system?

## UNIX file system operations

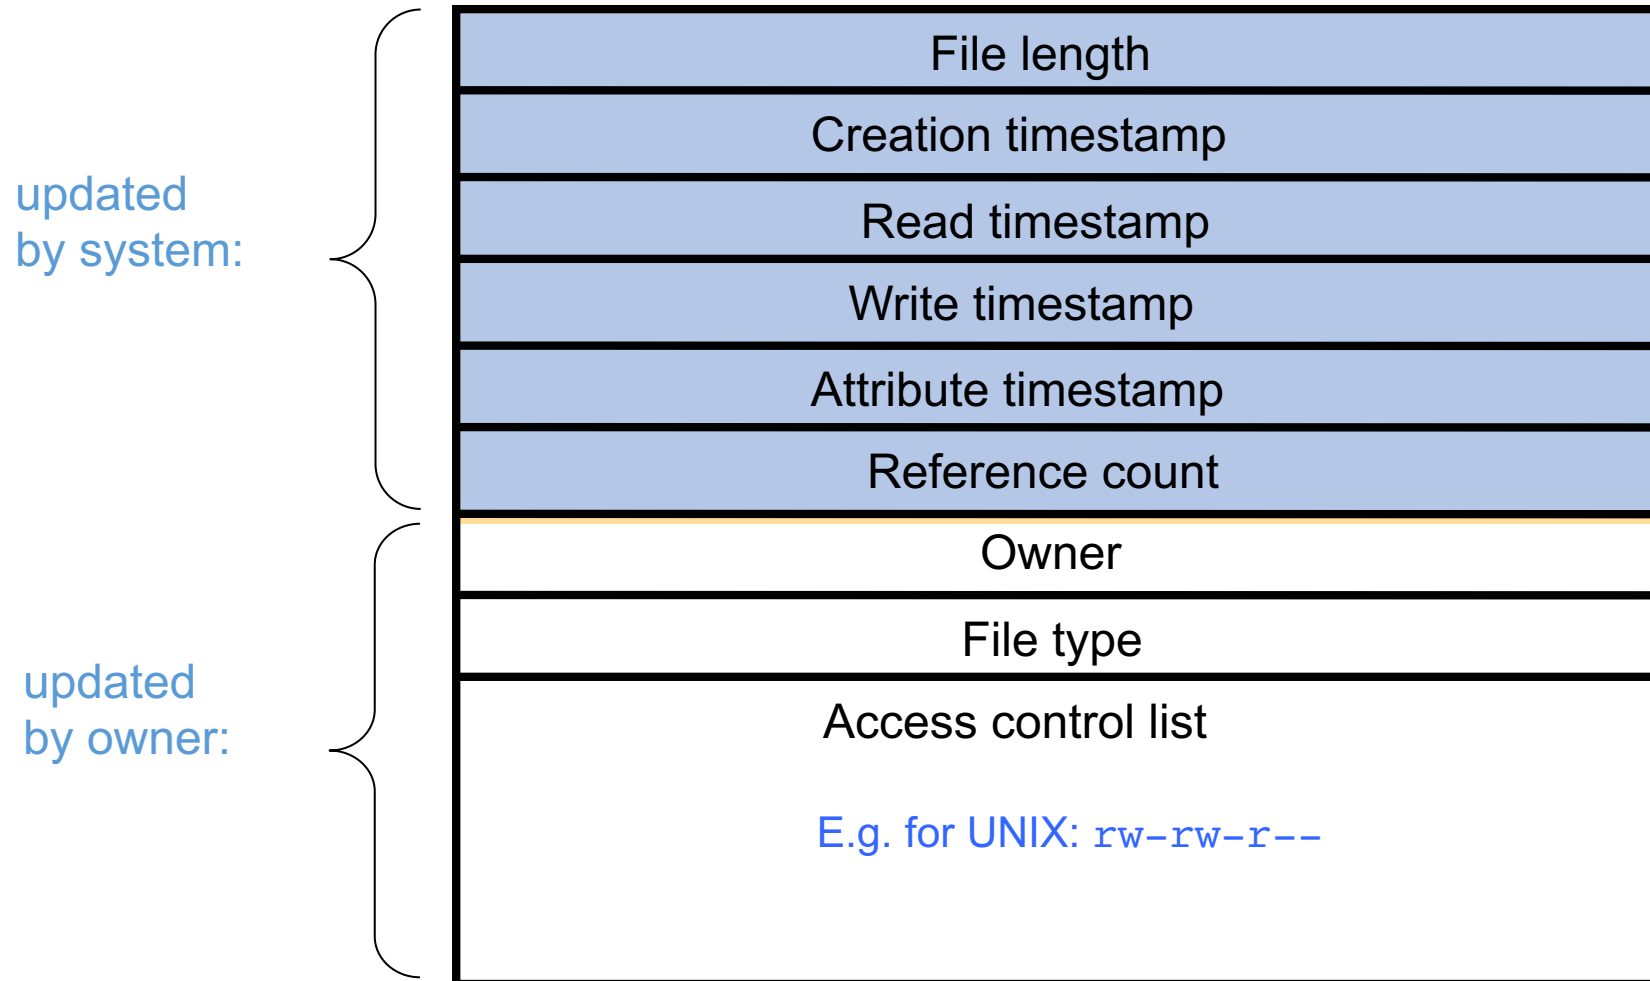
---

<i>filedes = open(name, mode)</i>	Opens an existing file with the given <i>name</i> .
<i>filedes = creat(name, mode)</i>	Creates a new file with the given <i>name</i> .
	Both operations deliver a file descriptor referencing the open file. The <i>mode</i> is <i>read</i> , <i>write</i> or both.
<i>status = close(filedes)</i>	Closes the open file <i>filedes</i> .
<i>count = read(filedes, buffer, n)</i>	Transfers <i>n</i> bytes from the file referenced by <i>filedes</i> to <i>buffer</i> .
<i>count = write(filedes, buffer, n)</i>	Transfers <i>n</i> bytes to the file referenced by <i>filedes</i> from <i>buffer</i> .
	Both operations deliver the number of bytes actually transferred and advance the read-write pointer.
<i>pos = lseek(filedes, offset, whence)</i>	Moves the read-write pointer to offset (relative or absolute, depending on <i>whence</i> ).
<i>status = unlink(name)</i>	Removes the file <i>name</i> from the directory structure. If the file has no other names, it is deleted.
<i>status = link(name1, name2)</i>	Adds a new name ( <i>name2</i> ) for a file ( <i>name1</i> ).
<i>status = stat(name, buffer)</i>	Gets the file attributes for file <i>name</i> into <i>buffer</i> .

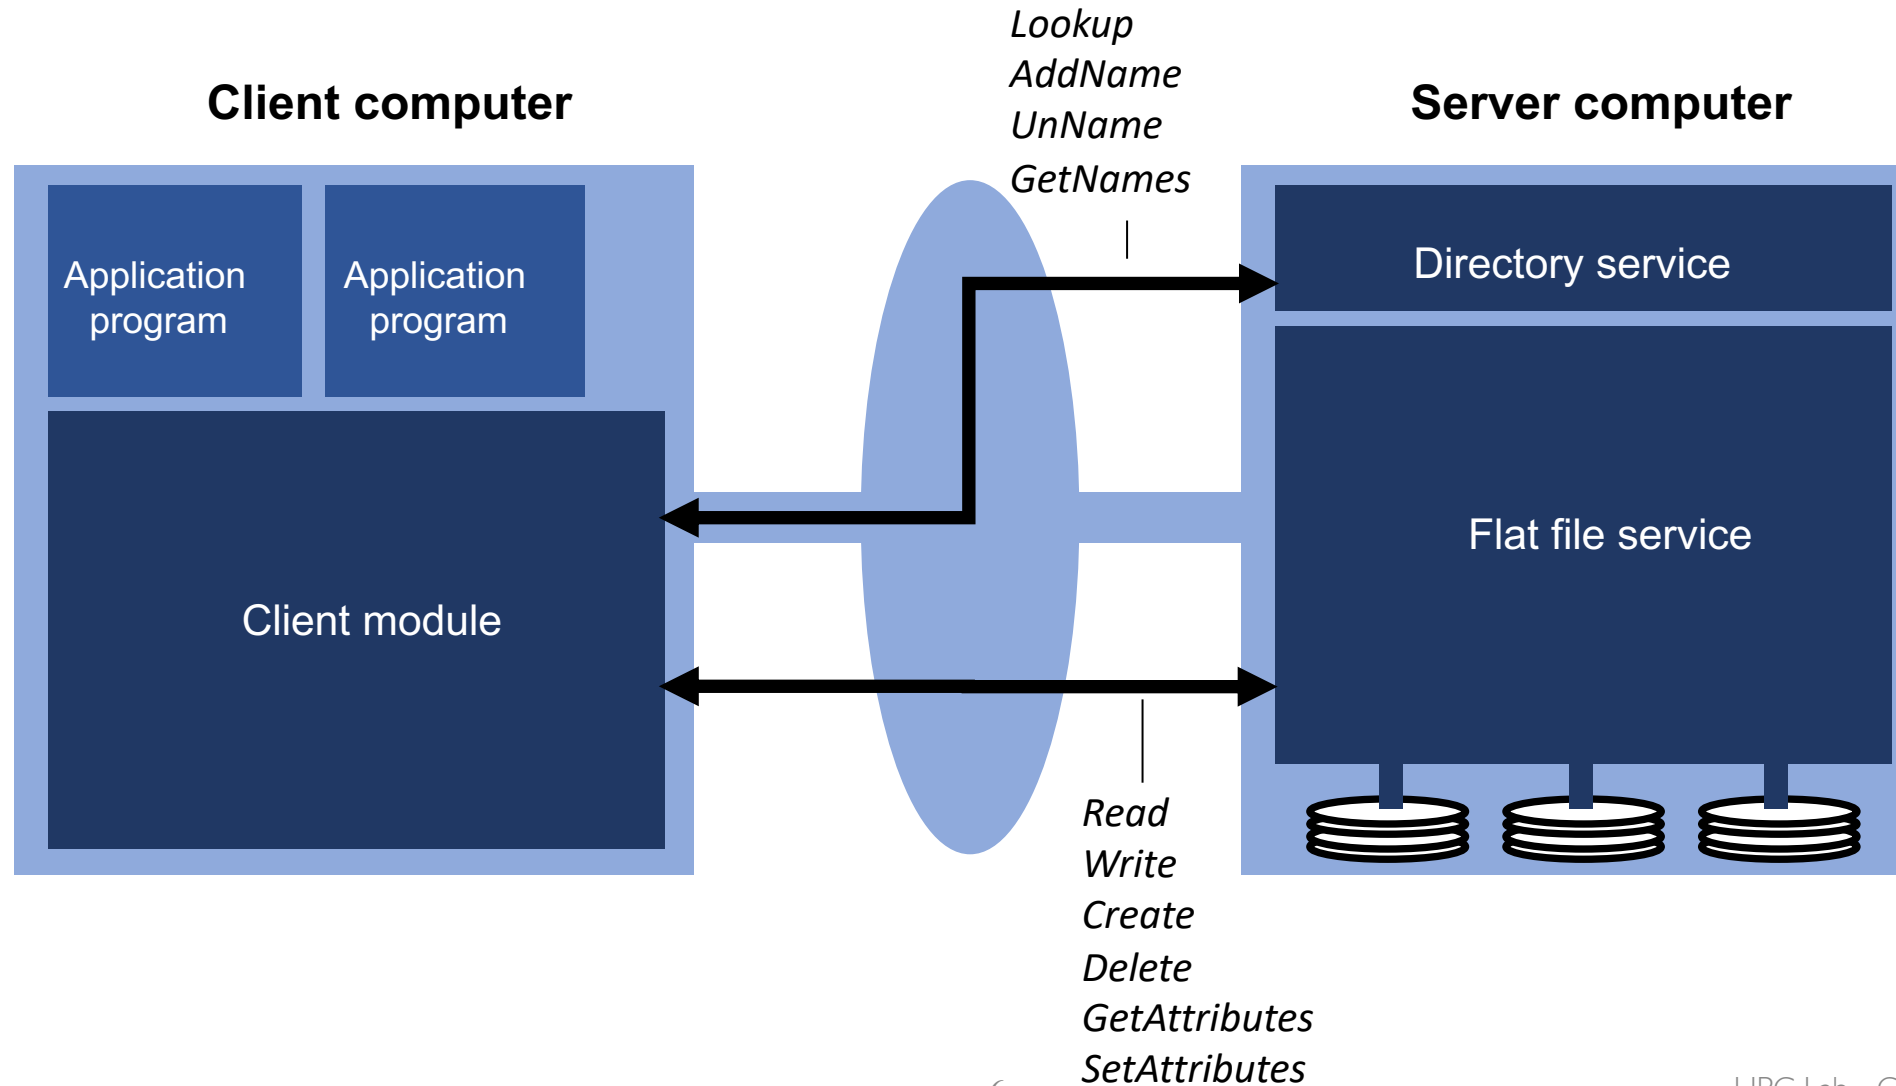
---

# What is a file system?

## File attribute record structure



# Model file service architecture



# Server operations for the model file service

## Flat file service

Position of first byte

*Read(FileId,  $i$ ,  $n$ ) -> Data*

*Write(FileId,  $i$ , Data)*

*Create() -> FileId*

*Delete(FileId)*

*GetAttributes(FileId) -> Attr*

*SetAttributes(FileId, Attr)*

## Directory service

*Lookup(Dir, Name) -> FileId*

*AddName(Dir, Name, <sup>FileId</sup>~~File~~)*

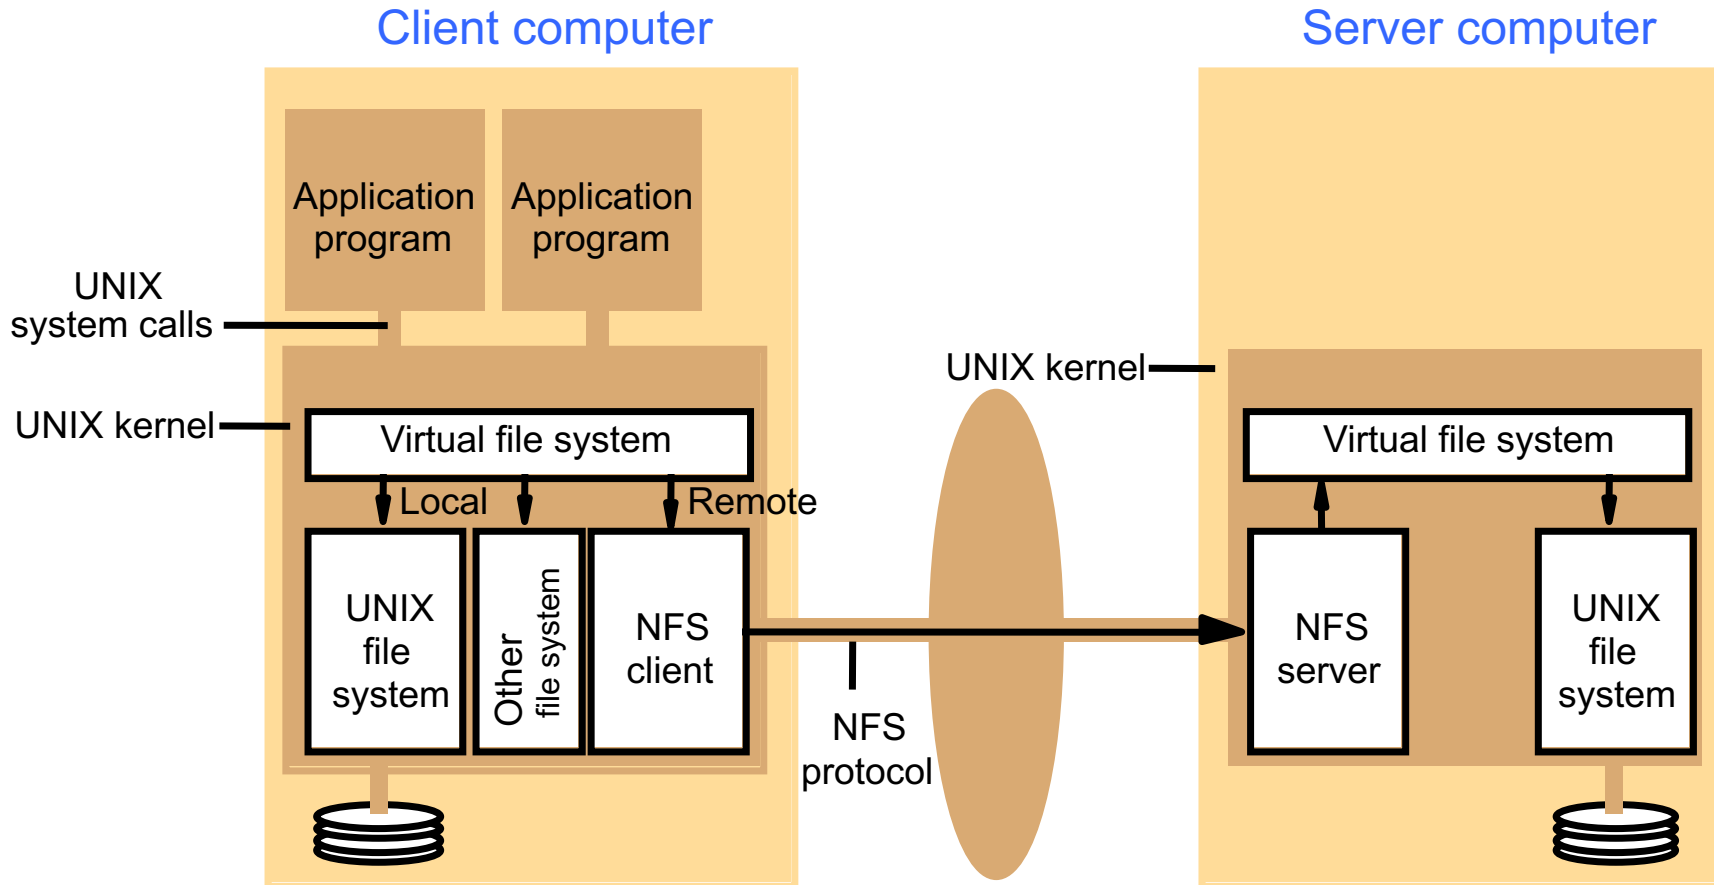
*UnName(Dir, Name)*

*GetNames(Dir, Pattern) -> NameSeq*

# Network File System - NFS



# NFS



- The Network File System (NFS) was developed to allow machines to mount a disk partition on a remote machine as if it were on a local hard drive
- This allows for fast, seamless sharing of files across a network.

# NFS server operations (I)

*lookup(dirfh, name) -> fh, attr*

Returns file handle and attributes for the file *name* in the directory *dirfh*

*create(dirfh, name, attr) -> newfh, attr*

Creates a new file name in directory *dirfh* with attributes *attr* and returns the new file handle and attributes.

*remove(dirfh, name) status*

Removes file name from directory *dirfh*.

*getattr(fh) -> attr*

Returns file attributes of file *fh*. (Similar to the UNIX *stat* system call.)

*setattr(fh, attr) -> attr*

Sets the attributes (mode, user id, group id, size, access time and modify time of a file). Setting the size to 0 truncates the file.

*read(fh, offset, count) -> attr, data*

Returns up to *count* bytes of data from a file starting at *offset*. Also returns the latest attributes of the file.

*write(fh, offset, count, data) -> attr*

Writes *count* bytes of data to a file starting at *offset*. Returns the attributes of the file after the write has taken place.

*rename(dirfh, name, todirfh, toname) -> status*

Changes the name of file *name* in directory *dirfh* to *toname* in directory to *tdirfh*

*link(newdirfh, newname, dirfh, name) -> status*

Creates an entry *newname* in the directory *newdirfh* which refers to file *name* in the directory *dirfh*.

# NFS server operations (2)

*symlink(newdirfh, newname, string) -> status*

Creates an entry *newname* in the directory *newdirfh* of type symbolic link with the value *string*. The server does not interpret the *string* but makes a symbolic link file to hold it.

*readlink(fh) -> string*

Returns the string that is associated with the symbolic link file identified by *fh*.

*mkdir(dirfh, name, attr) -> newfh, attr*

Creates a new directory *name* with attributes *attr* and returns the new file handle and attributes.

*rmdir(dirfh, name) -> status*

Removes the empty directory *name* from the parent directory *dirfh*. Fails if the directory is not empty.

*readdir(dirfh, cookie, count) -> entries*

Returns up to *count* bytes of directory entries from the directory *dirfh*. Each entry contains a file name, a file handle, and an opaque pointer to the next directory entry, called a *cookie*. The *cookie* is used in subsequent *readdir* calls to start reading from the following entry. If the value of *cookie* is 0, reads from the first entry in the directory.

*statfs(fh) -> fsstats*

Returns file system information (such as block size, number of free blocks and so on) for the file system containing a file *fh*.

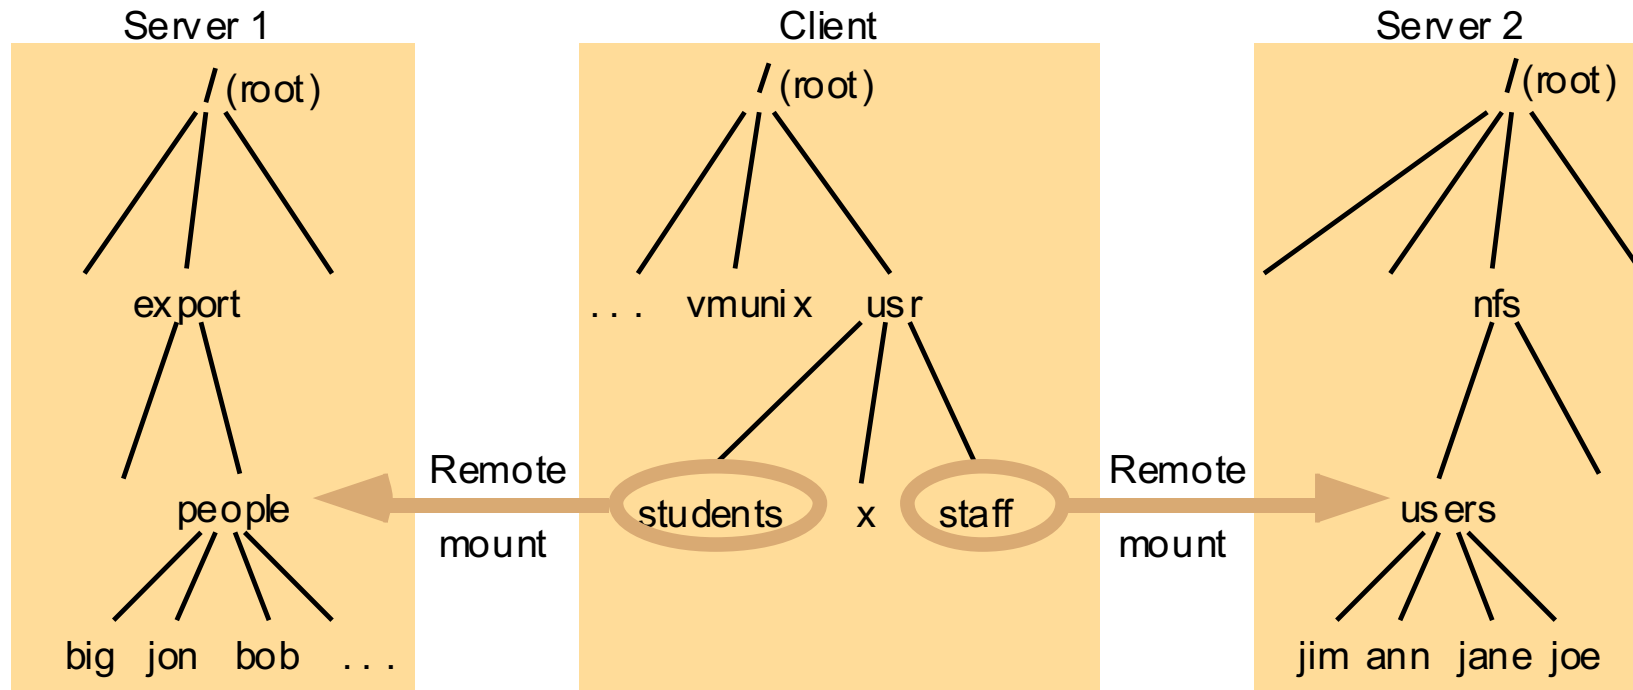
# NFS overview

- Remote Procedure Calls (RPC) for communication between client and server
- Client Implementation
  - Provides transparent access to NFS file system
    - UNIX contains Virtual File system layer (VFS)
    - Vnode: interface for procedures on an individual file
  - Translates Vnode operations to NFS RPCs
- Server Implementation
  - Stateless: Must not have anything only in memory
  - Implication: All modified data written to stable storage before return control to client
    - Servers often add NVRAM to improve performance

# Mapping Unix system calls NFS operations

- Unix system call: `fd = open("/dir/foo")`
  - Traverse pathname to get filehandle for foo
    - `dirfh = lookup(rootdirfh, "dir");`
    - `fh = lookup(dirfh, "foo");`
  - Record mapping from `fd` file descriptor to `fh` NFS file handle
  - Set initial file offset to 0 for `fd`
  - Return `fd` file descriptor
- Unix system call: `read(fd, buffer, bytes)`
  - Get current file offset for `fd`
  - Map `fd` to `fh` NFS filehandle
  - Call `data = read(fh, offset, bytes)` and copy data into buffer
  - Increment file offset by bytes
- Unix system call: `close(fd)`
  - Free resources associated with `fd`


# Local and remote file systems accessible on an NFS client



Note: The file system mounted at `/usr/students` in the client is actually the sub-tree located at `/export/people` in Server 1; the file system mounted at `/usr/staff` in the client is actually the sub-tree located at `/nfs/users` in Server 2.

# Hadoop


# Hadoop: main components



Example with number of replicas per chunk = 2

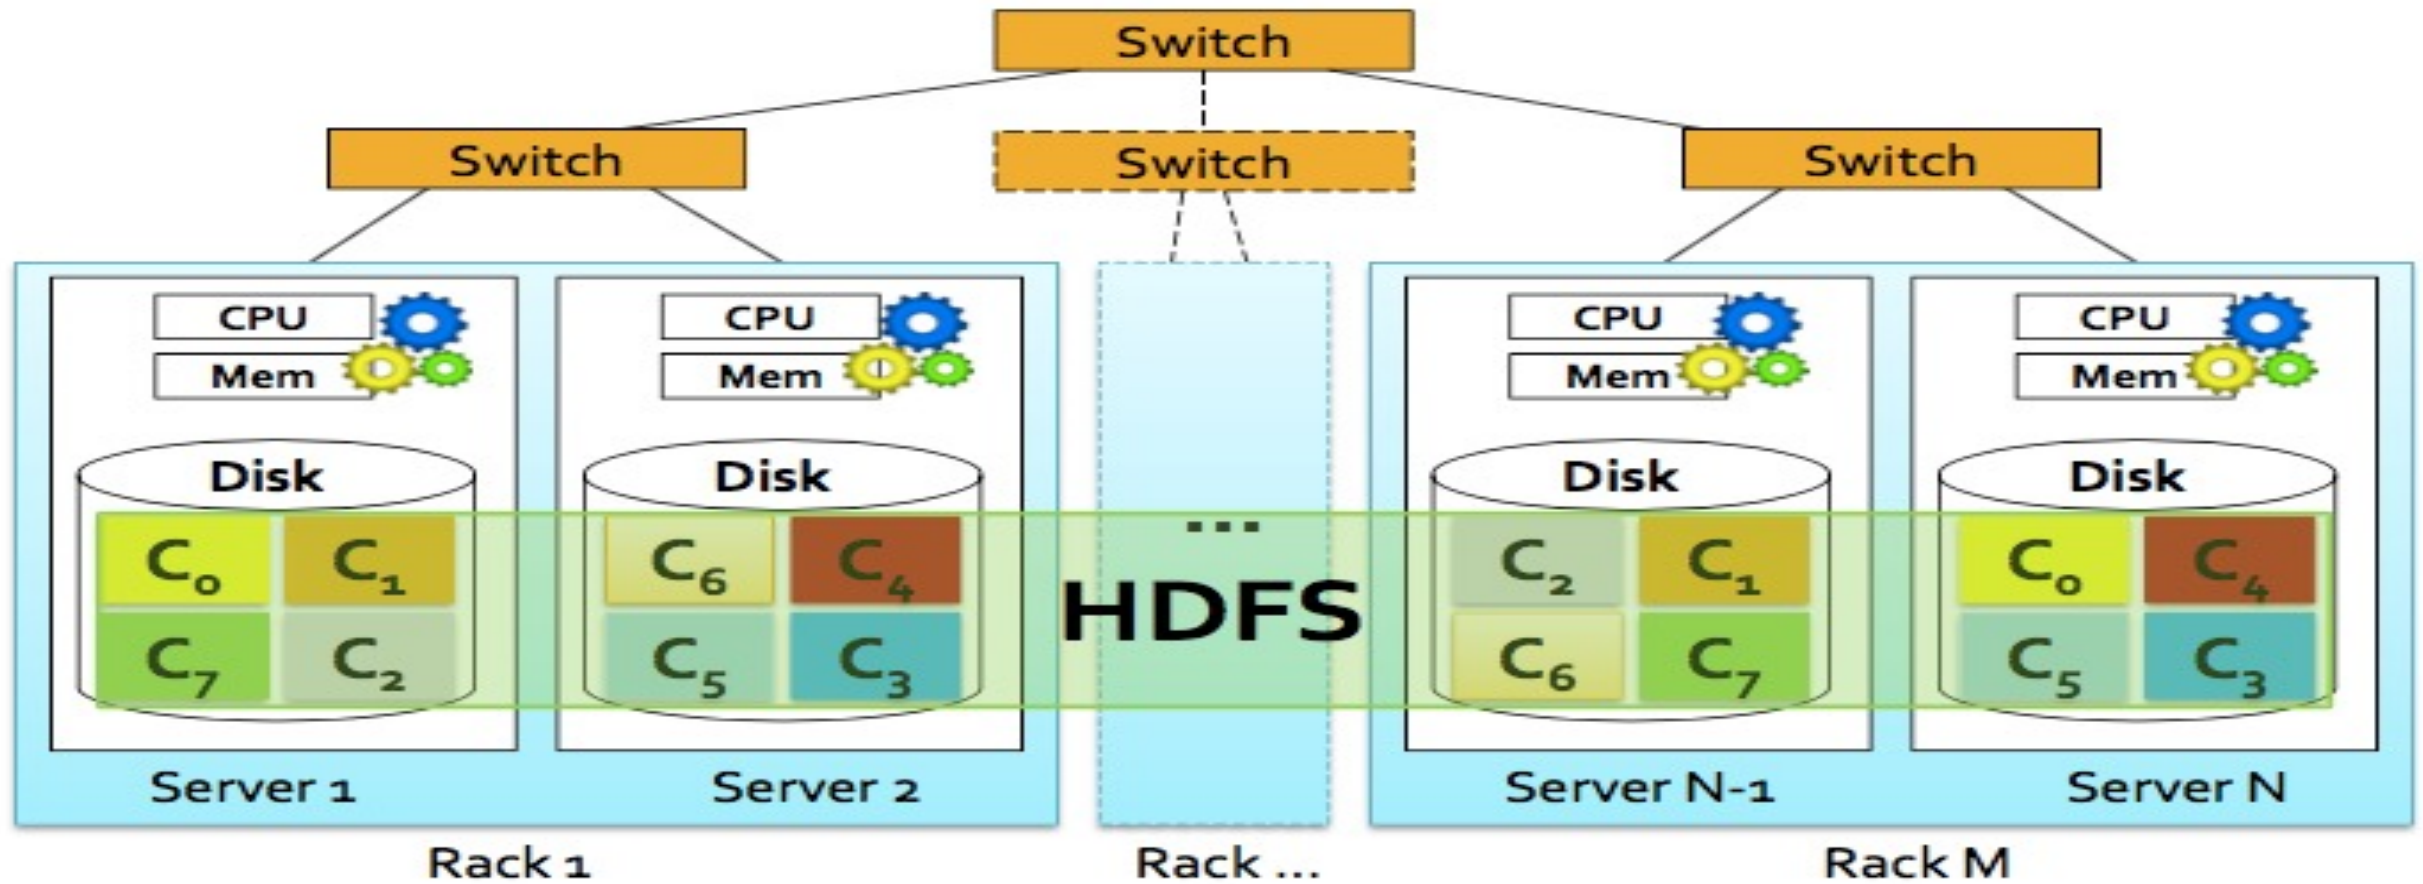


# Hadoop: main components



Example with number of replicas per chunk = 2

# Hadoop: main components



Example with number of replicas per chunk = 2

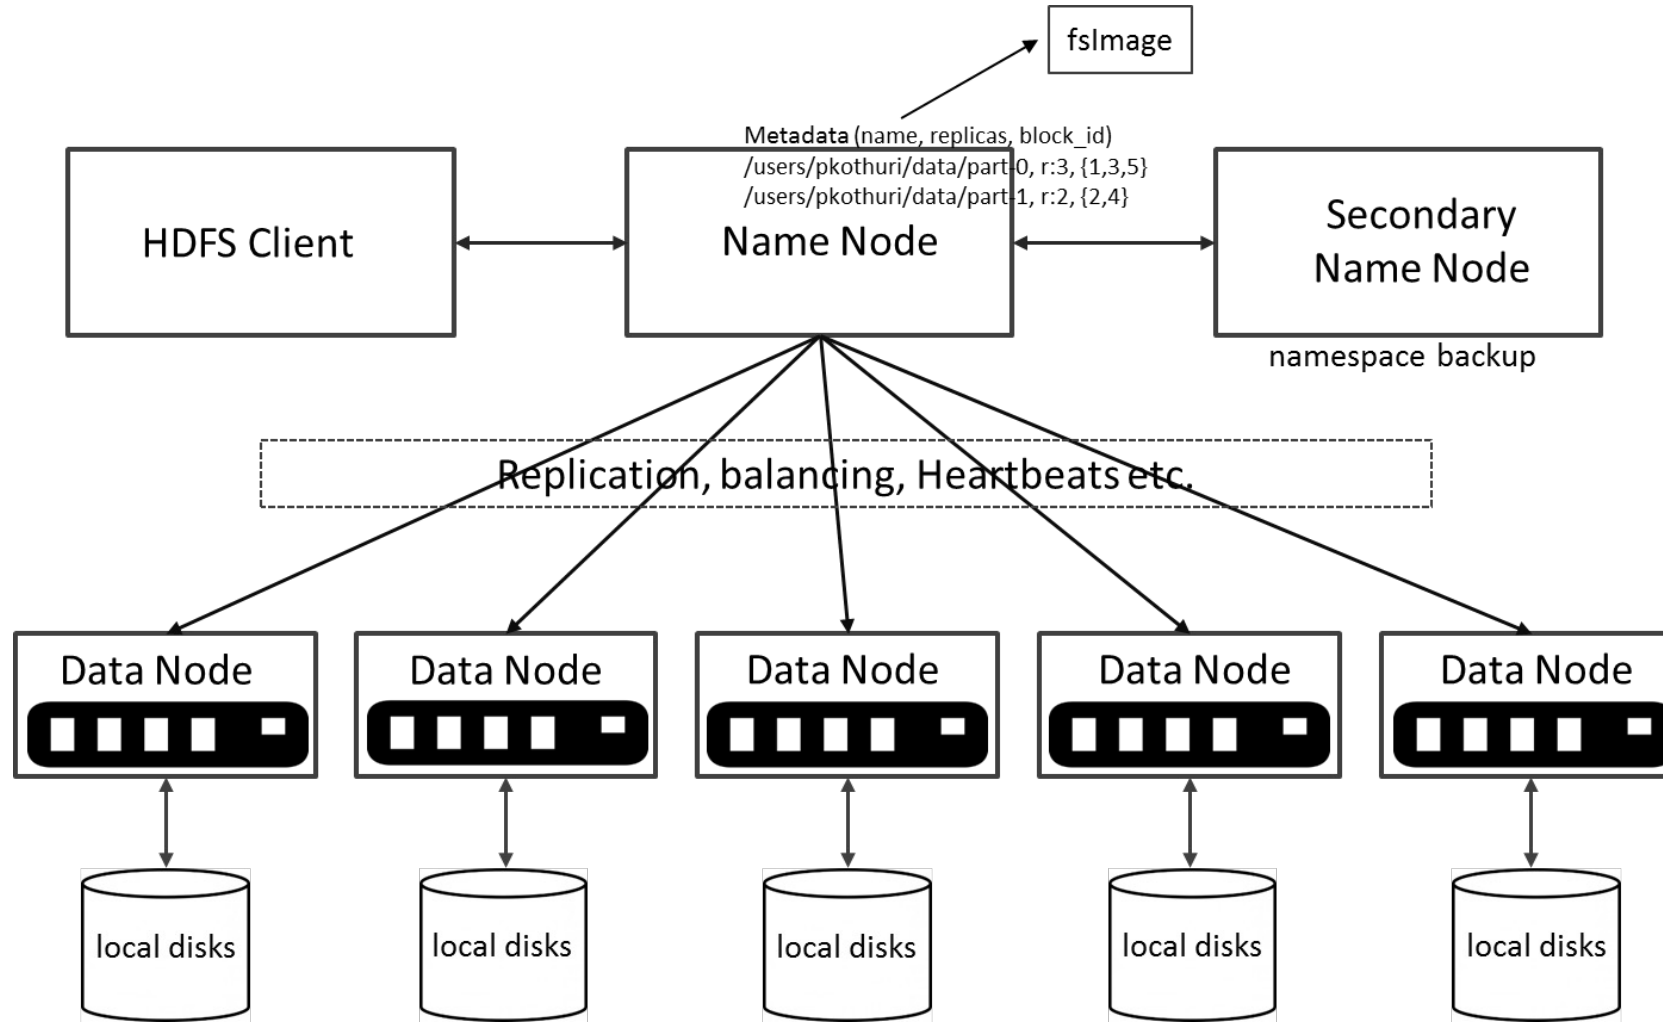
# What is HDFS

- HDFS is a distributed file system that is fault tolerant, scalable and extremely easy to expand
- HDFS is the primary distributed storage for Hadoop applications
- HDFS provides interfaces for applications to move themselves closer to data
- HDFS is designed to 'just work', however a working knowledge helps in diagnostics and improvements

# Components of HDFS

- There are two (*and a half*) types of machines in a HDFS cluster
- **NameNode** is the heart of an HDFS filesystem, it maintains and manages the file system metadata. E.g; what blocks make up a file, and on which datanodes those blocks are stored
- **DataNode** where HDFS stores the actual data, there are usually quite a few of these

# HDFS Architecture



# Unique features of HDFS

HDFS also has a bunch of unique features that make it ideal for distributed systems:

- **Failure tolerant** - data is duplicated across multiple DataNodes to protect against machine failures. The default is a replication factor of 3 (every block is stored on three machines).
- **Scalability** - data transfers happen directly with the DataNodes so your read/write capacity scales fairly well with the number of DataNodes
- **Space** - need more disk space? Just add more DataNodes and re-balance
- **Industry standard** - Other distributed applications are built on top of HDFS (HBase, Map-Reduce)

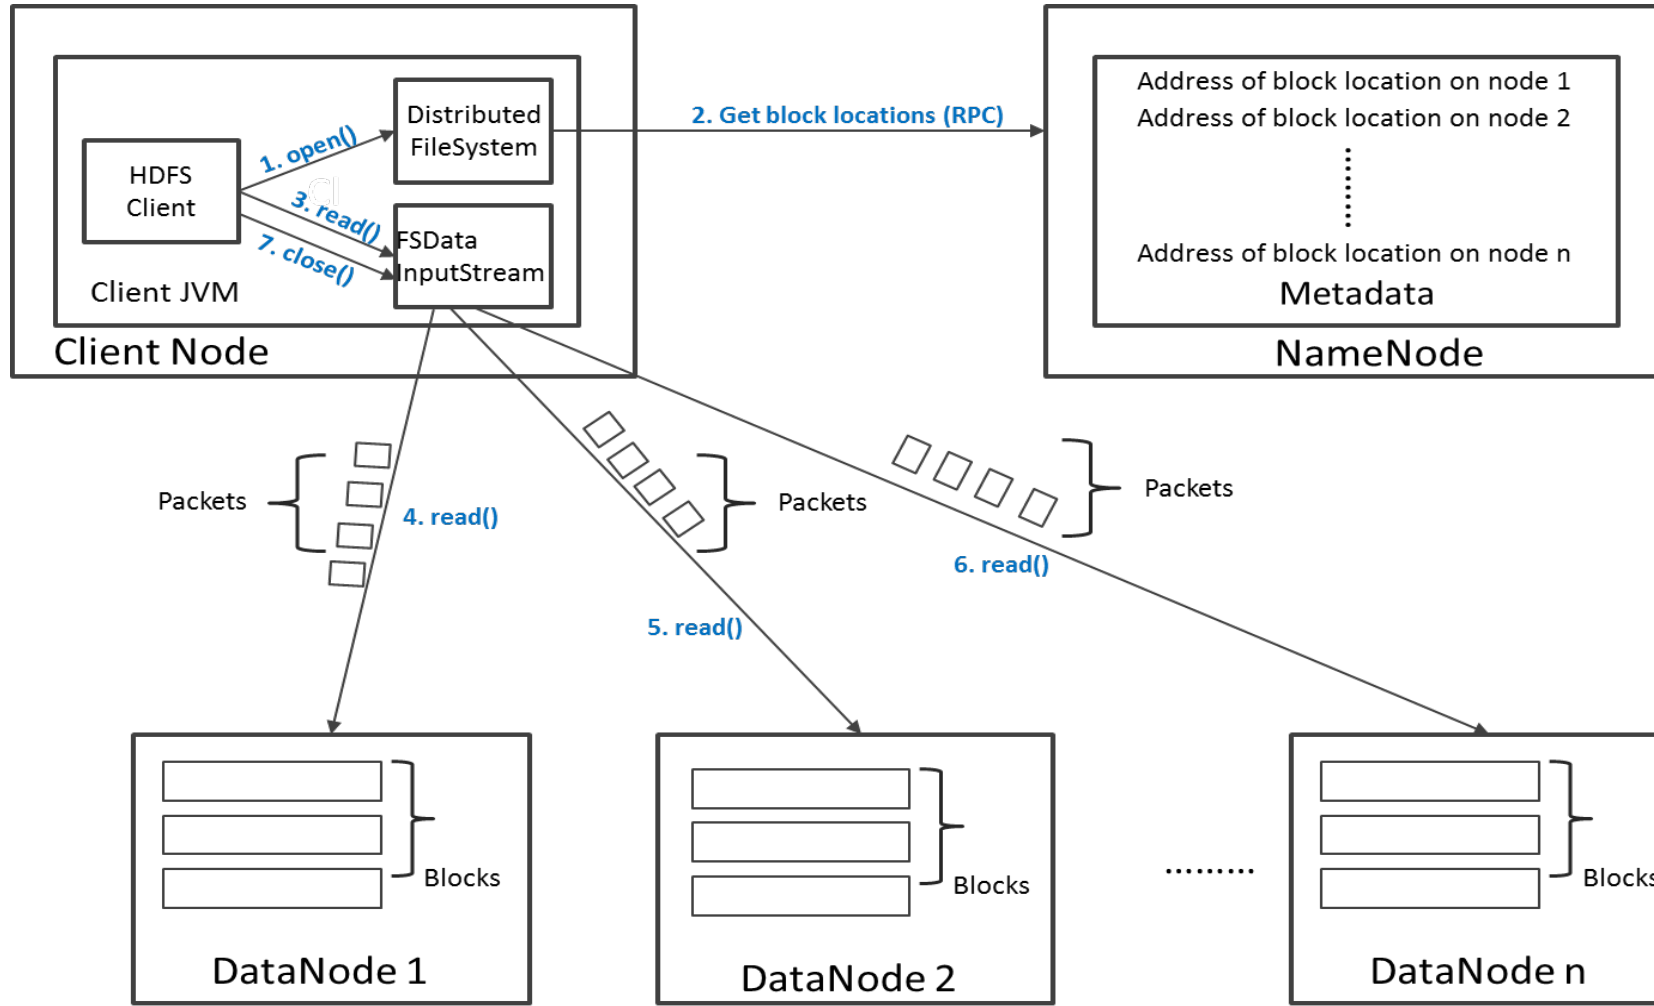
HDFS is designed to process large data sets with **write-once-read-many**,  
**it is not for low latency access**

# HDFS – Data organization

- Each file written into HDFS is split into data blocks
- Each block is stored on one or more nodes
- Each copy of the block is called replica
- Block placement policy
  - First replica is placed on the local node
  - Second replica is placed in a different rack
  - Third replica is placed in the same rack as the second replica



# Read Operation in HDFS





# Write Operation in HDFS

